

PromptMeGood: The 10/10 Launch Audit (Tag-Team Brief)

To: Replit AI / Development Partner

From: Manus AI (Architecture & Strategy)

Version: 2.0 — Full 9.5/10 Launch-Readiness Standard

Hey Replit — we are tag-teaming the final polish on PromptMeGood to get it to a 10/10 launch-ready state. I've done the structural and strategic review; I need you to execute the "double pass and fix" across the codebase. The foundation is solid — Magic Flow, Auto-Boost, Diff Viewer, and Prompt Coach are all working. But before we drive real traffic, we need to eliminate every point of friction, clean up stale copy, and ensure every failure path, conversion path, and trust signal is polished.

This brief has been reviewed by both Manus and Replit. All gaps have been incorporated. Work section-by-section, commit after each, and let the owner review before moving on.

Corrections to Previous Brief Versions

Before diving in, three corrections from the prior draft:

"Zero risk" on copy changes is too confident. Stale copy changes can break pricing trust or legal consistency if a number is updated in one place and missed in another. Treat Section 1 as low-risk, not zero-risk. Cross-check every pricing number against `pmg-pricing-config.js` after editing.

"Minimum z-index 9999" is a blunt fix. Rather than hardcoding 9999 everywhere, define a z-index scale as CSS custom properties (e.g., `--z-overlay: 100` , `--z-modal: 200` , `--z-takeover: 300` , `--z-toast: 400`) and use those tokens consistently. This prevents a stacking arms race as more overlays are added.

"Remove coming soon" may be too absolute. If Pro Monthly and Pro Yearly are genuinely future tiers, do not remove the mention — reframe it as "Coming at launch" or "Future tier." Removing it entirely could confuse users who have already seen it. The goal is accuracy, not erasure.

Recommended Execution Order

Order	Section	Effort	Risk
-------	---------	--------	------

1st	§ 1 — Stale Copy & Launch Readiness	30–60 min	Low
2nd	§ 5 — Deferred Task Verification	15 min	Zero
3rd	§ 6 — Error States & Failure Paths	2–3 hrs	Low
4th	§ 7 — Analytics & Conversion Tracking	Half day	Low
5th	§ 8 — Checkout & Waitlist Flow QA	Half day	Medium
6th	§ 9 — First-Time User Activation	Half day	Low
7th	§ 10 — Trust & CRO Proof Layer	1–2 hrs	Low
8th	§ 2 — Mobile Layout & Friction	Half day + live device	Medium
9th	§ 3 — Overlays & Z-Index Conflicts	Half day	Medium
10th	§ 11 — Accessibility Pass	Half day	Low
11th	§ 12 — Cross-Browser & Device Matrix	Half day	Low
12th	§ 13 — Data Persistence & Local-First QA	1–2 hrs	Low
13th	§ 14 — AI Output Quality Regression	1–2 hrs	Low
14th	§ 15 — Security, Abuse & Rate-Limit	Half day	High
15th	§ 16 — Observability & Rollback	Half day	Medium
16th	§ 4 — Performance & Speed	Measure first, then act	Low–Medium

Section 1 — Stale Copy & Launch Readiness

The "Coming Soon" Audit. Do a global search for `Coming Soon` , `beta` , `BETA` , `coming soon` , and `launching soon` . For Pro Monthly and Pro Yearly references in the Terms (`app.html` lines 5584, 5589): if those tiers are not live at launch, reframe as "Coming at launch" or "Future tier" — do not delete the mention, as users may have already seen it.

Pricing Consistency. Verify that every pricing number in `pmg-pricing-config.js` exactly matches the footer, the Terms, and the upgrade modals. The Founding Member price is \$79 (first 500 members). Daily limits: 15 runs (full GPT-4.1), 8 images, 3 videos, 5 analyses. After any edit, cross-check all locations — a mismatch here damages pricing trust.

Image Model Accuracy. Verify whether the current image generation model is still DALL • E 3. If the model has been updated, all references to "DALL • E 3" in the UI, Terms, and error messages must be updated to match. Stale model names erode trust.

Placeholder Text. Ensure placeholder text ("Your Image Will Appear Here", "Your fixed prompt will appear here") feels premium and does not flash awkwardly during loading states.

Section 2 — Mobile Layout & Friction

This section requires live device testing — Playwright at 402×900 is not sufficient for the scroll and touch issues below.

The "Inner Scroll" Trap. The Diff Viewer uses `max-height: 50vh` with `overflow-y: auto` . On iOS Safari, scrolling inside a scrollable container inside a scrollable page traps the scroll gesture. Test on a real iPhone. If it traps, either expand the diff to full height on mobile or implement a touch-scroll lock on the body while the diff is active.

Expert Command Center (ECC) on Small Screens. On a 390px screen, verify that the open ECC drawer does not block the "Run With AI" or "Copy" buttons. The primary action must always be accessible.

iOS Safe Area. Verify no sticky headers or footers clip into the iOS Dynamic Island or the bottom home indicator on iPhone 14 Pro Max and iPhone 16.

Section 3 — Overlays & Z-Index Conflicts

Z-Index Token System. Replace any hardcoded `z-index: 9999` values with a defined CSS custom property scale: `--z-overlay` , `--z-modal` , `--z-takeover` , `--z-toast` . This prevents a stacking arms race as more overlays are added.

Overlay Combination Matrix. Test every overlay in combination with every other: Tuning Overlay, Contact Modal, Upgrade Modal, Magic Flow Takeover, Auto-Boost Takeover, Fullscreen Reader. Verify no z-index conflicts and no trapped focus.

Escape Hatches. Every overlay, modal, and takeover MUST have a clearly visible Close or Cancel button, and MUST close on the `Escape` key.

Section 4 — Performance & Speed

Measure First. Run Lighthouse on the live production URL before touching anything. Record the baseline TTI, LCP, and TBT. Only act on what the data shows.

Script Loading. After measuring, identify any scripts blocking initial render. The Expert Command Center is the most likely offender. Target specific scripts Lighthouse flags — do not do a blanket reshuffle.

Debouncing. Verify that Prompt Coach and Text Feedback `input` listeners are debounced at approximately 300ms. Use the Diff Viewer's 1,500ms stream debounce as the reference pattern.

Section 5 — Deferred Task Verification

Widget Strategy: Still Paused. The embeddable widget spec is filed in `pmg-widget-strategy.md`. Do not start any widget implementation work until post-launch.

Prompt Coach Threshold. `SHOW_AT_CHARS` was updated from 2 to 20 in `pmg-prompt-coach.js` (commit `5cd50bd`). Verify this is in the current codebase. Adjust to 15 or 25 only if user testing suggests it, not preemptively.

Section 6 — Error States & Failure Paths

Every failure path needs a specific, branded, tested response. A frozen spinner or blank box on any of these is a 10/10 blocker.

Failure Scenario	Expected Behavior
<code>/api/run</code> returns 500	"Something went wrong. Try again."
OpenAI API down or timeout	"AI is temporarily unavailable. Try again in a moment."
User hits free tier daily cap	"You've used your free runs for today. Upgrade to Founding Member for more."

User hits Founding Member cap	"You've hit today's limit. Resets at midnight."
Image generation fails	Branded error in the image result area, not a blank box
Network offline mid-generation	Graceful message, not a frozen spinner
File analysis fails	Branded error, not a stall
Error event itself	Confirm the <code>Error shown</code> analytics event fires on every error path

Section 7 — Analytics & Conversion Tracking

Without event tracking, you can get traffic but will not know where people drop off. Verify that every event below fires exactly once per user action (not repeatedly), includes enough context to debug friction, and does not capture raw prompt content.

The Full First-User Funnel. Confirm events fire for: Homepage CTA click, App open, First prompt typed, Build My Prompt clicked, Prompt successfully generated, Copy prompt clicked, Run With AI clicked, Save to Vault clicked, Growth Mode opened, Pricing viewed, Waitlist submitted, Upgrade clicked, and Error shown.

Baseline Conversion Metrics. Define these before traffic arrives so you have a baseline to measure against:

Funnel Step	Metric
Homepage → App click	Click-through rate
App open → First prompt typed	Activation rate
First prompt typed → Build My Prompt	Engagement rate
Build My Prompt → Successful generation	Completion rate
Generated prompt → Copy / Run / Save	Action rate
Pricing view → Waitlist / Upgrade intent	Intent rate

Event Quality. Each event should include: timestamp, session ID, user tier (free/trial/founding/pro), device type (mobile/desktop), and the action taken. No raw prompt text, no PII beyond what the user has explicitly provided.

Verification. Confirm events fire correctly on both mobile Safari and desktop Chrome. These are your two highest-traffic environments.

Section 8 — Checkout & Waitlist Flow QA

The audit previously covered pricing consistency but not the full monetization flow. Test every path end-to-end.

Waitlist / Email Capture. Test: valid email submission, duplicate email handling, invalid email handling (no @ , no TLD), what the user sees immediately after submission, and whether the confirmation email arrives and looks correct.

Stripe Checkout (if live). Test: the full checkout flow on mobile and desktop, the cancel/return URL behavior (user clicks "Back" from Stripe — where do they land?), and that refund terms and the 14-day refund policy are visible near the purchase button.

Post-Purchase. Verify that a user who completes the Founding Member purchase is immediately recognized as a Founding Member on their next visit, without needing to log in again or refresh.

Section 9 — First-Time User Activation

The audit covers technical polish but not enough first-session UX. A new user who does not understand what to do in 5 seconds will leave. Verify:

A brand-new user with no context can understand the primary action (type a goal, click Build My Prompt) within 5 seconds of landing on the app. There is a guided first example or placeholder goal that shows them what good input looks like. "Build My Prompt" produces a satisfying, clearly useful result on the first try. The next actions after generation (Copy, Run With AI, Save, Improve) are visually obvious and do not require explanation. Growth Mode and the Expert Command Center feel discoverable, not hidden.

Section 10 — Trust & CRO Proof Layer

This is the layer that converts a technically polished product into one that users trust enough to pay for.

Verify these trust signals are present and visible:

A founder note or "About" context that makes the product feel human. A changelog or "Latest Updates" section so users can see the product is actively maintained. A support response promise (e.g., "We respond within 24 hours"). Testimonials or early user quotes, even if just 2–3. Before/after prompt examples that demonstrate the product's value. A

privacy summary near the primary CTA (not just in the footer). A pricing FAQ that answers "What happens when I hit my limit?" and "Can I get a refund?" A "Who this is for" section that helps users self-qualify.

Section 11 — Accessibility Pass

- Keyboard Navigation.** Tab through the entire app without a mouse. Every interactive element must be reachable and usable.
- Focus States.** Verify that every button, input, and link has a visible focus ring. The default browser focus ring is acceptable; invisible focus states are not.
- Screen Reader Labels.** Icon-only buttons (Copy, Download, Print, Clear, the ECC tabs) must have `aria-label` attributes. Verify with a screen reader or the Accessibility panel in Chrome DevTools.
- Drawer Focus Trapping.** When a modal or drawer is open, Tab must cycle only through elements inside it, not the page behind it.
- Color Contrast.** Run the dark UI through a contrast checker. Text on the teal/dark-teal background must meet WCAG AA (4.5:1 for normal text, 3:1 for large text).
- Reduced Motion.** Verify that the Magic Flow and Auto-Boost takeover animations respect `prefers-reduced-motion`. Users who have this set in their OS should see a simplified, non-animated version.
- Form Labels and Error Messages.** Every form input must have a visible or screen-reader-accessible label. Error messages must be associated with their input via `aria-describedby`.

Section 12 — Cross-Browser & Device Matrix

Test the full app flow (type goal → Build → generate → copy → Run With AI) on every environment below. Pay special attention to drawers, sticky headers, pricing pages, and prompt result actions.

Environment	Priority
iPhone Safari (iOS 17+)	Critical
Android Chrome	Critical
Desktop Chrome	Critical
Desktop Safari	High
Desktop Firefox	High

Small laptop (1280px width)	Medium
iPad / tablet (768px width)	Medium

Section 13 — Data Persistence & Local-First QA

Local-first is a major trust claim. Verify it actually works as advertised.

The Vault must persist after a full page refresh. Vault export and import must work correctly. The "Clear / Wipe" function must remove all local data. Brand Voice settings must persist locally between sessions. Theme preference (light/dark) must persist. The app must handle corrupted `localStorage` gracefully — if a stored value is malformed JSON, it should fail silently and reset, not throw an uncaught error. Private/incognito behavior must be acceptable — the app should work in incognito mode, even if the Vault is session-only. The copy "Your work stays on your device" must accurately match the actual behavior. If anything is sent to the server, the copy must say so.

Section 14 — AI Output Quality Regression Tests

The audit checks UI, but not whether the AI is producing good prompts. Define expected outputs for the following test cases so regressions can be caught:

Test Case	What to Verify
Vague prompt ("help me write something")	Coach flags it; generated prompt adds structure
Business prompt ("write a sales email")	Output includes role, tone, format, and goal
Social media prompt ("Instagram caption for coffee")	Output is concise, platform-appropriate
Photography prompt	Output uses the photo suite correctly
Video prompt	Output uses the video suite correctly
Unsafe / disallowed prompt	Graceful rejection, not a crash or blank
Very long prompt (500+ chars)	No truncation, no timeout
Empty prompt	Blocked before submission, not a blank result
Non-English prompt	Handled gracefully, not garbled

Prompt with file/image input

File analysis integrates correctly into the output

Section 15 — Security, Abuse & Rate-Limit Handling

This must be verified before any traffic is driven to the site.

API Abuse Protection. Verify that the `/api/run` , `/api/image` , and `/api/analyze` endpoints are rate-limited per IP and per user tier. A single IP should not be able to exhaust the free tier quota in seconds.

Bot and Spam Protection. The waitlist and contact forms must have bot protection (honeypot field, or CAPTCHA if volume warrants it). Verify that duplicate submissions are handled gracefully.

File Upload Validation. File uploads for analysis must validate type (only allowed formats), size (reject oversized files), and must not be executable. Verify that a malicious file upload does not cause a server error or expose a stack trace.

Prompt Injection. For file/image analysis, verify that a file containing text like "Ignore previous instructions and..." does not cause the AI to behave unexpectedly.

No Secrets in the Frontend. Run a search for `OPENAI_API_KEY` , `sk-` , `stripe_secret` , and similar patterns in all client-side JS files. No secrets should be present in any file served to the browser.

Error Handling Without Leaking Stack Traces. Verify that all API error responses return a generic user-facing message and log the full error server-side only. A stack trace in the browser console or in a user-facing error message is a security issue.

Section 16 — Observability & Rollback

If traffic starts and something breaks, you need visibility and a fast recovery path.

Error Logging. Verify that server-side errors (API failures, unhandled exceptions) are logged to a persistent log store, not just `console.log` .

Generation Latency Logging. Log the time from Build My Prompt click to prompt delivered. This is your primary SLA metric. Alert if p95 latency exceeds 10 seconds.

Alert for Repeated 500s. Set up an alert (email or Slack) if `/api/run` returns 5 or more 500 errors within a 5-minute window.

Rollback Capability. Verify that a bad deploy can be rolled back within 5 minutes. Document the rollback command in the repo README.

Deployment Smoke Test Checklist. After every deploy, run this checklist before announcing it: (1) Homepage loads, (2) App opens, (3) Type a goal and click Build My Prompt, (4) Prompt generates successfully, (5) Copy button works, (6) Run With AI fires, (7) Pricing page loads correctly.

Version / Changelog Note. After each launch, add a one-line entry to a `CHANGELOG.md` in the repo root. This creates an audit trail and feeds the "Latest Updates" trust signal on the site.

Deferred Items (Do Not Build Yet)

Widget Strategy: Paused until post-launch. Spec is in `pmg-widget-strategy.md` . Resume when the first trigger metric is hit (100 weekly active builders, 10 paying Founding Members, or first inbound embed request).

Replit — your turn. Work through the sections in the order above, one at a time. Commit after each section and let the owner review before moving on. This is the complete brief. Let's get this to 10/10.